

CSE 847 – HW 4 – Andrew Haas

GitHub link: <https://github.com/AndrewEHaas/cse847hw4>

First, I implemented the following python code to drive my regression:

```
# ==== Additional Methods ==== #

sigmoid = lambda data : 1/(1 + np.exp(-1*data))
predict = lambda data, weights : np.where(sigmoid(data @ weights) > .5, 1, 0)

def loss(data: Array, weights: Array, labels: Array) -> float:
    s = sigmoid(data @ weights)
    return (-1/data.shape[0])*np.sum(labels*np.log(s) + (1-labels)*np.log(1-s))

def gradient(data: Array, weights: Array, labels: Array) -> Array:
    s = sigmoid(data @ weights)
    g = np.zeros(weights.shape)
    for i in range(data.shape[0]):
        g += np.reshape((labels[i,:] - s[i,:])*data[i,:], g.shape)
    return (-1/data.shape[0])*g

def l1_loss(data: Array, weights: Array, labels: Array, par: float) -> float:
    return loss(data, weights, labels) + par*np.linalg.norm(weights, ord=1)/data.shape[0]

def l1_gradient(data: Array, weights: Array, labels: Array, par: float) -> float:
    return gradient(data, weights, labels) + par*np.sign(weights)/data.shape[0]

# ===== #
```

These methods were then called by the following functions. First, standard logistic regression (0,1 encoding):

```
def logistic_train(data: Array, labels: Array, lr: float, epsilon: float = 1e-5, maxiter: int = 1000) ->
Array:
    """
    parameters:
        data : n * (d + 1) with n samples and d features, where col d+1 is all ones (corresponding to
intercept)

        labels : n * 1 vector of class labels (0 or 1)

        lr : learning rate

        epsilon: optional parameter specifying convergence - if the change in absolute difference in
predictions, from one iteration to the next
when averaged across input features is less than epsilon, halt (default 1e-5)

        maxiter: optional parameter that specifies the maximum number of iterations (default 1000)

    return: 1 * (d + 1) vector of weights corresponding to columns of data - last entry is bias term
    """
    # instantiate weight vector, store previous weight iteration
    weights = np.zeros((data.shape[1],1))
    prev = np.copy(weights)

    #ls = [] # list to store loss values at each iteration

    for iter in range(maxiter):

        # update weights
        weights -= lr*gradient(data, weights, labels)

        # compute and store loss
        #ls.append(loss(data, weights, labels))

        # check break condition, halting of necessary
        if np.linalg.norm(prev - weights, ord=2) < epsilon:
            return weights
        prev = np.copy(weights)

    #plt.plot(range(maxiter), ls, label='learning_rate: ' + str(lr))
    #plt.xlabel('iteration')
    #plt.ylabel('loss function')
    #plt.legend(loc='best')
    #plt.savefig('test.jpg')
    return weights
```

(note: the plotting code shown here is commented, though used in one test that will be shown shortly)

Then, l_1 regularized logistic regression, performed via sub gradient descent:

```
def logistic_l1_train(data: Array, labels: Array, par: float, lr: float, epsilon: float = 1e-5, maxiter:
int = 5000):
    """
    parameters:
        data : n * (d + 1) with n samples and d features, where col d+1 is all ones (corresponding to
intercept)

        labels : n * 1 vector of class labels (0 or 1)

        par : regularization parameter

        lr : learning rate

        epsilon: optional parameter specifying convergence - if the change in absolute difference in
predictions, from one iteration to the next
                when averaged across input features is less than epsilon, halt (default 1e-5)

        maxiter: optional parameter that specifies the maximum number of iterations (default 1000)

    return: 1 * (d + 1) vector of weights corresponding to columns of data - last entry is bias term
    """
    # instantiate weight vector, store previous weight iteration
    weights = np.zeros((data.shape[1],1))
    prev = np.copy(weights)

    #ls = [] # list to store loss values

    for iter in range(maxiter):

        # update weights
        weights -= lr*l1_gradient(data, weights, labels, par)

        # compute and store loss
        #ls.append(l1_loss(data, weights, labels, par))

        # check break condition, halting if necessary
        if np.linalg.norm(prev - weights, ord=2) < epsilon:
            return weights
        prev = np.copy(weights)

    #plt.plot(range(maxiter), ls, label='par: ' + str(par))
    #plt.xlabel('iteration')
    #plt.ylabel('loss function')
    #plt.legend(loc='best')
    #plt.savefig('test.jpg')
    return weights
```

Using this code, I implemented the following tests.

Problem 1:

First, I had to determine whether my code even worked. To do this, I made convergence plots on the first dataset using both types of regression, and a wide variety of parameters. The following code was used to drive the test:

```
def test1a():
    # simply looks at convergence for various parameters of dataset 1

    # read data
    data = np.loadtxt('data_spam.txt')
    labels = np.loadtxt('labels_spam.txt')


    # append column of 1 ones for bias term
    data = np.concatenate((data, np.ones((data.shape[0],1))), axis=1)


    # reshape labels
    labels = np.reshape(labels, (4601, 1))

    plt.figure(figsize=(16,13), dpi=100)
    for lr in [.1, .01, .001, .0001]:
        logistic_train(data, labels, lr)

    plt.figure(figsize=(16,13), dpi=100)
    for par in [0, .01, .1, .2, .3, .4, .5, .6, .7, .8, .9, 1]:
        logistic_l1_train(data, labels, par, .01, maxiter=5000)
```

Producing the following output, and verifying convergence:





Next, I performed the test specified in part one of the assignment – splitting the spam dataset into a testing and training set and uses varying portions of the training set to test our algorithm.

```
def test1b():
    # tests accuracy with various training sets

    # read data
    data = np.loadtxt('data_spam.txt')
    labels = np.loadtxt('labels_spam.txt')

    # append column of 1 ones for bias term
    data = np.concatenate((data, np.ones((data.shape[0],1))), axis=1)

    # reshape labels
    labels = np.reshape(labels, (4601, 1))

    # split data and labels for testing
    x_test, y_test = data[2000:, :], labels[2000:, :]

    table = pt.PrettyTable()
    table.field_names = ['Training Size', 'Testing Misclassifications']

    for training_size in [200, 500, 800, 1000, 1500, 2000]:
        w = logistic_train(data[:training_size:], labels[:training_size], .05, maxiter=7500)
        prediction = predict(x_test, w)
        temp = 0
        for i in range(y_test.shape[0]):
            if prediction[i,:] != y_test[i,:]:
                temp += 1
        table.add_row([training_size, temp])

    print(table)
```

Which generated the following output (notice this test was performed without regularization, with a learning rate of .05 and 7500 iterations:

```
+-----+-----+
| Training Size | Testing Misclassifications |
+-----+-----+
|      200     |             198           |
|      500     |             179           |
|      800     |             175           |
|     1000     |             175           |
|     1500     |             164           |
|     2000     |             159           |
+-----+-----+
```

Problem 2:

The following code was used to test problem two. I considered a coefficient zero if it were less than .0001, and used a step size of .01.

```
def test2():
    # read data
    data = loadmat('ad_data.mat')
    x_train, y_train = data['X_train'], data['y_train']
    x_test, y_test = data['X_test'], data['y_test']

    # change encoding
    y_train = np.where(y_train == 1, 1, 0)
    y_test = np.where(y_test == 1, 1, 0)

    table = pt.PrettyTable()
    table.field_names = ['Reg Param', 'Testing Misclassifications', 'Nonzero Weights']
    for par in [0, .01, .1, .2, .3, .4, .5, .6, .7, .8, .9, 1]:
        w = logistic_l1_train(x_train, y_train, par, .01, epsilon=1e-20, maxiter=5000)
        prediction = predict(x_test, w)

        # count misclassifications
        temp1 = 0
        for i in range(y_test.shape[0]):
            if prediction[i,:] != y_test[i,:]:
                temp1 += 1
        # count nonzero weights
        temp2 = 0
        for i in range(w.shape[0]):
            if abs(w[i,:]) > 1e-3:
                temp2 += 1

        table.add_row([par, temp1, temp2])

    print(table)
```

Which produced the following table, showing behavior much as one would expect as the regularization parameter grew:

Reg Param	Testing Misclassifications	Nonzero Weights
0	19	116
0.01	19	116
0.1	20	105
0.2	19	86
0.3	19	75
0.4	19	68
0.5	19	60
0.6	19	48
0.7	19	45
0.8	19	41
0.9	19	38
1	19	31